

Learning Turbulence Closures without Losing the Physics

RANS, inverse PINN, neural regression, and reproducible evidence

FlowMLLab | Week 14 | Ehsan Roohi | Graduate research lecture

1. The question: better k, not just a prettier velocity curve

Research foundation: Lars Davidson's pyCALC-RANS and PINN-NN turbulence-model workflow. DNS reference: Myoungkyu Lee and Robert D. Moser. These lecture explanations and newly generated figures are a FlowMLLab teaching adaptation, not Davidson's original slides.

A mean velocity profile can look convincing while the predicted turbulent kinetic energy is substantially wrong. We will distinguish the observable we want to improve from the equations that constrain it, reproduce source-code operations, and ask what survives feedback through a CFD solver.

Learning outcomes: explain a closure coefficient; reconstruct an inverse-physics target; train the original small neural regression model; separate interpolation from transfer; and report both successful checks and failed convergence criteria.

Suggested format: 75-minute lecture plus a 90-minute computer lab. Prerequisites: Reynolds averaging, k-omega modeling, finite-volume residuals, neural regression and automatic differentiation.

2. What does a turbulence closure actually change?

Reynolds averaging introduces the unresolved stress tensor. An eddy-viscosity closure relates its deviatoric part to the resolved strain. This makes the mean-flow equations solvable, but it does not make the turbulence model exact.

For a steady, fully developed channel, the mean flow depends strongly on the eddy viscosity $\nu_t = k/\omega$. Increasing k alone increases ν_t , which changes the mean velocity. A coordinated change to ω can preserve the ratio while allowing a different turbulent-energy level.

$$\begin{aligned}\nu_t &= k/\omega, & P_k &= \nu_t (dU/dy)^2 \\ 0 &= P_k - 0.09C_k k\omega + \frac{d}{dy}[(\nu + \nu_t/\sigma_k)dk/dy] \\ 0 &= \frac{5}{9}\frac{\omega}{k}P_k - C_{\omega 2}\omega^2 + \frac{d}{dy}[(\nu + \nu_t/\sigma_\omega)d\omega/dy]\end{aligned}\tag{1}$$

The three learned quantities have different jobs: σ_k controls turbulent diffusion of k ; C_k scales its destruction; $C_{\omega 2}$ controls ω destruction. $C_{\omega 2}$ is an absolute coefficient here, not another multiplicative correction to 0.075. The unmodified values are $\sigma_k = 2$, $C_k = 1$ and $C_{\omega 2} = 0.075$.

The corrected channel solver also ties σ_ω to σ_k through $\min(2\sigma_k, 2)$; it is not an independently trained fourth network.

Think before fitting: if the velocity already agrees well but k does not, which equation term should you inspect first? Why is fitting U alone insufficient to identify all three correction functions?

3. Establish the right baseline




Figure 1. Classical and table-based PINN profiles against Lee-Moser DNS. Markers are fresh restarted solves, not a verified reproduction of the final PINN-NN curves in Figure 8 of the paper.

Use the actual unmodified channel case, not a table that happens to include RANS in its filename. The supplied notebook's embedded profile matches the upstream PINN-corrected directory byte for byte. Calling that profile the classical baseline hides much of the model discrepancy.

The figures distinguish archived source data from freshly executed results. A restarted simulation is a new numerical execution, but not a cold-start demonstration. Record the grid, initial fields, stopping criterion and achieved residual alongside every accuracy claim.

These percentages describe the distributed table-based PINN stage, not a verified reproduction of the final PINN-NN curves in the paper's Figure 8. Section 5.1 reports good mean-velocity predictions and improved k for the final model. Do not attribute our particular velocity-error increase to that final-model claim.

All profile norms must identify their weighting. An unweighted Euclidean norm on a stretched grid emphasizes regions with many stored points. A cell-width-weighted norm answers a different question. Neither is an experimental uncertainty estimate.

4. One research pipeline, several evidence levels




Figure 2. Four stages of the attributed research workflow, each requiring a different kind of evidence.

The inverse problem supplies a spatial diffusion correction using reference statistics. Supervised networks then express the correction coefficients through local flow features. Deployment evaluates those networks inside the iterative RANS calculation.

Do not collapse these stages into one accuracy number. A low PINN residual tests the inverse equation; a low NN regression error tests coefficient approximation; a converged CFD residual

tests the numerical solve; agreement with DNS tests selected physical observables.

Each stage can succeed while a later stage fails. A model fitted on a known channel profile can have an excellent random-point test score without establishing transfer to separation, a new geometry or a new Reynolds number.

5. The inverse-physics step

Treat the turbulent diffusion coefficient as an unknown function $a(y)$. With reference production, dissipation and k , the channel energy balance gives a differential constraint on a . A neural representation allows its derivative to be calculated by automatic differentiation.

$$r(y) = a'(y)k'(y) + (\nu + a(y))k''(y) + P_k(y) - \varepsilon(y)$$

$$\mathcal{L} = \sum_i r(y_i)^2 + 1000\mathcal{L}_{BC} \quad (2)$$

$$\sigma_k(y) = \nu_t(y)/a(y)$$

The released implementation combines a summed interior residual with a boundary penalty of 1000. Those numerical weights are implementation choices, not universal physical constants. Derivatives of reference k are calculated numerically before interpolation onto the RANS grid.

Near a stationary point of k , its first derivative is small; inference of a from the balance becomes sensitive. Reference discretization, endpoint conditions, positivity, smoothing and the exact transport terms therefore matter. Small training loss alone cannot certify a uniquely identified physical closure.

The long PINN optimization and archived checkpoints must be labeled separately from a fresh supervised fit. Loading coefficients previously obtained by PINN is not training a PINN during this lab.

Our full 200,000-epoch source run reached a minimum loss near 1.24 but finished near 231. The original script does not select the best checkpoint. Do not substitute the minimum for the final result or claim reproduction of the paper's reported final loss of 4.

6. Three targets and their units




Figure 3. Released spatial correction coefficients and unmodified model constants.

The channel scripts use $\delta = u_{\tau} = 1$ and $\nu = 1/\text{Re}_{\tau}$. Thus stored U and k coincide numerically with U^+ and k^+ , but stored ω uses the outer time scale δ/u_{τ} . To obtain ω^+ multiply stored ω by ν/u_{τ} squared.

At deployment the inputs are $\nu_t/(y^*u_{\tau})$ and total shear magnitude divided by u_{τ} squared. In the training table, $u_{\tau} = 1$ simplifies these expressions. The released c_k feature calculation caps total shear at 0.995.

The balance-derived targets include caps and local smoothing. These choices must be recorded, not described as exact DNS truth. The k correction and ω correction compensate for changing the diffusion model; they are not independent measurements.

Important execution finding: the released balance script does not regenerate the bundled targets from the bundled inputs. Maximum absolute differences are about 0.525 for the two-column C_k table, 0.0198 for $C_{\omega 2}$ and 0.299 for σ_k . We preserve both versions and do not claim end-to-end target reproduction.

Lab check: confirm matching coordinates, finite values, target shape and file hashes before training. A matching array length alone does not prove that two profiles share the same grid.

7. Reproduce the network before changing it




Figure 4. Original c_k regression protocol with an explicit initialization seed. Random-point accuracy is not whole-case transfer.

The released c_k script uses raw local features with MinMax scaling, two hidden layers of ten ReLU units, a linear scalar output and mean-squared error. Training uses SGD, learning rate 0.04, batch size one and 1000 epochs. Its random 80/20 point split uses `random_state = 42`.

For repeatability the FlowMLLab runner fixes the PyTorch initialization seed and uses one CPU thread. These are explicit orchestration additions. The source creates a learning-rate scheduler but does not step it; reproducing the script therefore means retaining the constant rate.

The original scaler sees the complete profile before the split. Preserve that behavior only in the clearly labeled reproduction. In a new predictive experiment fit preprocessing on development data alone. Never silently improve the protocol and still call the result an exact reproduction.

The historical notebook instead used log-transformed features, standardization, tanh and L-BFGS. That is a legitimate comparison model, but not the original algorithm.

8. The interpolation control




Figure 5. PCHIP and the separately labeled tanh-MLP control over the historically inspected held-out interval.

A contiguous gap in one profile asks whether a model can interpolate across missing wall-normal samples. Compare against shape-preserving PCHIP in $\log(y^+)$, not only against a constant predictor. PCHIP uses position directly and is intentionally a strong profile-specific control, not a transferable local closure.

The interval $100 \leq y^+ \leq 400$ is already part of the inspected teaching material. It is held out from fitting, but no longer an unopened prospective test. Do not rename it blind to make the evidence sound stronger.

Report coefficient errors and the frozen-field destruction diagnostic separately. Multiplying coefficient errors by $0.09 \cdot k \cdot \omega$ changes their physical weighting, but does not account for the response of k and ω to a changed closure.

If interpolation wins, retain that result. The educational objective is sound inference, not ensuring that the neural network wins every comparison.

9. Put the closure back into the solver




Figure 6. Source-solver residuals. NN5200 is an assembled case; neither NN run meets the requested $1e-6$ criterion.

In a coupled run, the current fields produce the features, the networks produce coefficients, and the discretized momentum and turbulence equations update the fields. Clipping and relaxation can help stabilize this feedback, but also change the deployed map.

Check finite fields, positivity, residual trends, wall shear, profile changes and coefficient bounds. For the unit-pressure-gradient half-channel, the expected wall shear is one in the selected normalization. This is an independent physical check, not a substitute for all equation residuals.

The source package includes corrected initialization code but a stale preassembled executable. Rebuilding from the corrected source is essential. The exact source hashes and any compatibility changes belong in the run record.

A residual below an operational tolerance must not be described as satisfying a stricter original tolerance. Equally, a tiny residual does not establish mesh independence or correctness of the turbulence model.

10. A reproducibility ledger

Table 1. Actual iterations and original stopping gates.

Fresh run	Iterations	Residual	Original gate
baseline	20000	3.12e-13	NOT MET
pinn	2	9.31e-07	PASS
nn10000	40000	1.57e-05	NOT MET
nn5200	40000	2.39e-03	NOT MET




Figure 7. Fresh 200,000-epoch checkpoint continuation and the released diffusion profile. Minimum and final loss are distinct.

For each result keep: source URL and archive hash; input hashes; software versions; seed; grid; restart or cold-start status; actual iterations and stopping criterion; coefficient clipping; and the reference used for each error measure.

The supplied notebook is retained as a historical input and executed separately. The corrected teaching notebook does not overwrite it. Source code, retained author checkpoints, newly trained models and new solver outputs are separate categories.

The instructor reported Davidson's permission in the shared conversation. That statement is not a general license to relicense or distribute the complete source archive. The external code stays in the local download area; cite Davidson for the solver and workflow, and Lee-Moser for DNS.

The corresponding 2026 journal paper is the principal research reading. The older ETMM15 paper is useful historical context, but should not be the only reference for the later PINN-NN pipeline.

11. What this lecture does not prove

Reproduction of a known channel result is not a new discovery and not a prospective generalization study. A restarted run does not establish robustness from arbitrary initial conditions. One mesh does not establish discretization independence.

Changing Reynolds number while using a source-provided trained model can test execution of a published transfer case, but it is not a newly blinded benchmark. Keep the training case, model-development history and evaluation case explicit.

A research extension would freeze the complete preprocessing and model-selection protocol, reserve unopened whole cases, include a classical closure baseline, report failures, and examine uncertainty and out-of-support features. Separated flows require additional evidence; do not infer cavity or airfoil validity from channel interpolation.

Discussion: which is more informative for deployment, a two-percent coefficient error or a stable solve with slightly improved k but degraded wall shear? Explain why the answer depends on the intended use.

12. Assessment and primary references

Student submission: one provenance table; the true-baseline and corrected-flow comparison; original-network and interpolation-control metrics; a convergence and wall-shear check; and a 200-word claim boundary. Identify one implementation detail that would materially change the interpretation if omitted.

Extension exercise: derive the integrated total-shear balance for the half-channel. Compare its prediction with the numerical profile. Explain why differentiation errors and cell-face interpolation can affect a pointwise diagnostic even when the finite-volume balance converges.

L. Davidson, Using Physics Informed Neural Network (PINN) and Neural Network (NN) to Improve a k - ω Turbulence Model, *Journal of Turbulence* 27(7), 187-208 (2026). DOI: 10.1080/14685248.2026.2665148. Public preprint: arxiv.org/abs/2511.12493v3.

L. Davidson, pyCALC-RANS: A Python Code for Two-Dimensional Turbulent Steady Flow (2026). Report and source distribution: cfd-sweden.se/lada/pyCALC-RANS.html. The PINN-NN distribution page records the September 2026 initialization fixes.

M. Lee and R. D. Moser, Direct numerical simulation of turbulent channel flow up to Re_τ approximately 5200, *Journal of Fluid Mechanics* 774, 395-415 (2015). DOI: 10.1017/jfm.2015.268.